

Data stream
SEA / RBF / RADAR

Clean or six poisoned conditions
Label / concept / splice; 15% or 25%

Predict before
learning

Drift monitor
(when used)

Fixed response
No / immediate /
confirmed reset

Train candidate
learner

Reset arms: ADWIN, HDDM-W, Hellinger or D3-inspired

No-reset control: continue learning without a response monitor

Unscreened
Accept the candidate directly

Screened: protected validation
Accept if accuracy does not fall;
otherwise retain the active model

Compare matched prediction, poison admission and legitimate withholding.