

Attacker access: offline labelled source archive, including future rows
Fixed interventions; no model queries or detector-state access

Label poisoning



$(x, y) \rightarrow (x, \text{altered label})$

Features remain unchanged.

Concept replacement



$(x, y) \rightarrow (\text{generated } x, \text{generated } y)$

250-row intervals; marginal sampling and a projection rule.

Splice copying



$(x, y) \rightarrow (\text{copied } x, \text{copied } y)$

250-row intervals copied from another source segment.

Protected boundary: original validation rows and evaluation truth
The attacker cannot alter released protected feedback or scoring references.