

Shared inputs: SEA, RBF and RADAR
One clean + six poisoned conditions; same learner and reserved rows



Experiment 1: fixed responses
Compare no reset, immediate reset and confirmed reset.
Repeat each comparison with and without update screening.



Experiment 2: learned reset control
Replace the fixed rule with two frozen Q-learning policies.
Keep ADWIN and protected screening unchanged.



Experiment 3: observation filtering
Filter before monitoring and training; retain validation.
Compare filters, monitors and matched no-reset controls.

Evaluate prediction, poison admission, legitimate withholding and response activity.